

```

1 from google.colab import files
2 import pandas as pd
3
4 # Upload the file
5 uploaded = files.upload()

```



Choose Files Titanic-Dataset.csv

- Titanic-Dataset.csv(text/csv) - 61194 bytes, last modified: 12/24/2021 - 100% done
Saving Titanic-Dataset.csv to Titanic-Dataset.csv

```

1 from sklearn.model_selection import train_test_split
2 from sklearn.preprocessing import StandardScaler
3 from sklearn.linear_model import LogisticRegression
4 from sklearn.metrics import accuracy_score, classification_report

```

```

1 titanic_df = pd.read_csv('Titanic-Dataset.csv')
2
3 # Display the first few rows of the dataset
4 titanic_df.head()

```



	PassengerId	Survived	Pclass	Name	Sex	Age	SibSp	Parch	Ticket	
0	1	0	3	Braund, Mr. Owen Harris	male	22.0	1	0	A/5 21171	7.
1	2	1	1	Cumings, Mrs. John Bradley (Florence Briggs)	female	38.0	1	0	PC 17599	71.

Next steps:

[Generate code with titanic_df](#)



[View recommended plots](#)

[New interactive sheet](#)

```
1 titanic_df.info()
```



```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 891 entries, 0 to 890
Data columns (total 12 columns):
#   Column      Non-Null Count  Dtype
---  -
0   PassengerId  891 non-null    int64
1   Survived     891 non-null    int64
2   Pclass       891 non-null    int64
3   Name         891 non-null    object
4   Sex          891 non-null    object
5   Age          714 non-null    float64
6   SibSp        891 non-null    int64
7   Parch        891 non-null    int64
8   Ticket       891 non-null    object
9   Fare         891 non-null    float64
10  Cabin        204 non-null    object
11  Embarked     889 non-null    object

```

```
dtypes: float64(2), int64(5), object(5)
memory usage: 83.7+ KB
```

```
1 # Drop columns that won't be used in the model
2 titanic_df.drop(['PassengerId', 'Name', 'Ticket', 'Cabin'], axis=1, inplace=True)
3
4 # Handle missing values
5 titanic_df['Age'].fillna(titanic_df['Age'].median(), inplace=True)
6 titanic_df['Embarked'].fillna(titanic_df['Embarked'].mode()[0], inplace=True)
7
8 # Convert categorical variables to numeric
9 titanic_df = pd.get_dummies(titanic_df, columns=['Sex', 'Embarked'], drop_first=True)
10
11 # Display the processed dataframe
12 print(titanic_df.head())
```

```

Survived  Pclass  Age  SibSp  Parch  Fare  Sex_male  Embarked_Q  \
0         0      3  22.0      1      0   7.2500      True      False
1         1      1  38.0      1      0  71.2833     False      False
2         1      3  26.0      0      0   7.9250     False      False
3         1      1  35.0      1      0  53.1000     False      False
4         0      3  35.0      0      0   8.0500      True      False

Embarked_S
0         True
1        False
2         True
3         True
4         True
```

```
1 # Define features (X) and target (y)
2 X = titanic_df.drop('Survived', axis=1)
3 y = titanic_df['Survived']
4
5 # Split the data into training and testing sets
6 X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)
7
8 # Standardize the features
9 scaler = StandardScaler()
10 X_train = scaler.fit_transform(X_train)
11 X_test = scaler.transform(X_test)
```

```
1 # Initialize the model
2 model = LogisticRegression()
3
4 # Train the model
5 model.fit(X_train, y_train)
```

```

LogisticRegression
LogisticRegression()
```

```

1 # Predict on the test set
2 y_pred = model.predict(X_test)
3
4 # Evaluate the model
5 accuracy = accuracy_score(y_test, y_pred)
6 report = classification_report(y_test, y_pred)
7
8 print(f"Accuracy: {accuracy}")
9 print("Classification Report:\n", report)

```



Accuracy: 0.8100558659217877

Classification Report:

	precision	recall	f1-score	support
0	0.83	0.86	0.84	105
1	0.79	0.74	0.76	74
accuracy			0.81	179
macro avg	0.81	0.80	0.80	179
weighted avg	0.81	0.81	0.81	179

```


1 import matplotlib.pyplot as plt
2 import seaborn as sns
3 from sklearn.metrics import roc_curve, roc_auc_score, confusion_matrix
4
5 # Set the style for seaborn
6 sns.set(style="whitegrid")

```

```

1 # Plotting the correlation heatmap
2 plt.figure(figsize=(10, 8))
3 sns.heatmap(titanic_df.corr(), annot=True, fmt=".2f", cmap='coolwarm')
4 plt.title("Correlation Heatmap")
5 plt.show()


```




```

1 # Plotting the survival count
2 sns.countplot(x='Survived', data=titanic_df)
3 plt.title("Count of Survival")
4 plt.show()

```



```
1 from sklearn.ensemble import RandomForestClassifier
2
3 # Train a Random Forest model
4 rf_model = RandomForestClassifier(random_state=42)
5 rf_model.fit(X_train, y_train)
6
7 # Get feature importances
8 importances = rf_model.feature_importances_
9 feature_names = X.columns
10
11 # Plot feature importances
12 plt.figure(figsize=(10, 6))
13 sns.barplot(x=importances, y=feature_names)
14 plt.title("Feature Importance")
15 plt.show()
```



```

1 # Calculate ROC curve and AUC
2 y_pred_prob = model.predict_proba(X_test)[: , 1]
3 fpr, tpr, thresholds = roc_curve(y_test, y_pred_prob)
4 roc_auc = roc_auc_score(y_test, y_pred_prob)
5
6 # Plot ROC curve
7 plt.figure(figsize=(8, 6))
8 plt.plot(fpr, tpr, label=f"ROC Curve (area = {roc_auc:.2f})")
9 plt.plot([0, 1], [0, 1], 'k--')
10 plt.xlabel("False Positive Rate")
11 plt.ylabel("True Positive Rate")
12 plt.title("ROC Curve")
13 plt.legend(loc="lower right")
14 plt.show()

```

